

Package ‘bartMachine’

December 5, 2013

Type Package

Title Bayesian Additive Regression Trees

Version 1.0

Date 2013-12-2

Author Adam Kapelner and Justin Bleich

Maintainer Adam Kapelner <kapelner@wharton.upenn.edu>

Description An advanced implementation of Bayesian Additive Regression Trees (BART, Chipman et al, 2010)

License MIT

Depends R (>= 2.14.0), rJava (>= 0.8-4), car, randomForest, missForest

SystemRequirements Java (>= 1.6.27)

R topics documented:

bart_machine_get_posterior	2
bart_machine_num_cores	3
bart_predict_for_test_data	4
build_bart_machine	6
build_bart_machine_cv	10
calc_credible_intervals	11
calc_prediction_intervals	12
check_bart_error_assumptions	14
cov_importance_test	15
destroy_bart_machine	17
dummify_data	18
get_sigsq	19
get_var_counts_over_chain	20
get_var_props_over_chain	21
init_java_for_bart_machine_with_mem_in_mb	22
interaction_investigator	23
investigate_var_importance	25
k_fold_cv	26
pd_plot	28

plot_convergence_diagnostics	30
plot_y_vs_yhat	31
predict.bartMachine	32
print.bartMachine	34
rmse_by_num_trees	35
set_bart_machine_num_cores	36
summary.bartMachine	37
var_selection_by_permute_response_cv	38
var_selection_by_permute_response_three_methods	40

Index	43
--------------	-----------

bart_machine_get_posterior

Get Full Posterior Distribution

Description

Generates draws from posterior distribution of $\hat{f}(x)$ for a specified set of observations.

Usage

```
bart_machine_get_posterior(bart_machine, new_data)
```

Arguments

bart_machine	An object of class “bartMachine”.
new_data	A data frame containing observations at which draws from posterior distribution of $\hat{f}(x)$ are to be obtained.

Value

Returns a list with the following components:

y_hat	Posterior mean estimates. For regression, the estimates have the same units as the response. For classification, the estimates are probabilities.
new_data	The data frame with rows at which the posterior draws are to be generated. Column names should match that of the training data.
y_hat_posterior_samples	The full set of posterior samples of size num_iterations_after_burn_in for each observation. For regression, the estimates have the same units as the response. For classification, the estimates are probabilities.

Note

This function is parallelized by the number of cores set in [set_bart_machine_num_cores](#).

Author(s)

Adam Kapelner and Justin Bleich

See Also

[calc_credible_intervals](#), [calc_prediction_intervals](#)

Examples

```
## Not run:
#Regression example

#generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)

#get posterior distribution
posterior = bart_machine_get_posterior(bart_machine, X)
print(posterior$y_hat)

#destroy BART model
destroy_bart_machine(bart_machine)

#Classification example

#get data and only use 2 factors
data(iris)
iris2 = iris[51:150,]
iris2$Species = factor(iris2$Species)

#build BART classification model
bart_machine = build_bart_machine(iris2[,1 : 4], iris2$Species)

#get posterior distribution
posterior = bart_machine_get_posterior(bart_machine, iris2[,1 : 4])
print(posterior$y_hat)

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

bart_machine_num_cores

Get Number of Cores Used by BART

Description

Returns number of cores used by BART

Usage

```
bart_machine_num_cores()
```

Details

Returns the number of cores currently being used by parallelized BART functions

Value

Number of cores currently being used by parallelized BART functions.

Author(s)

Adam Kapelner and Justin Bleich

See Also

[set_bart_machine_num_cores](#)

Examples

```
bart_machine_num_cores()
```

```
bart_predict_for_test_data
```

Predict for Test Data with Known Outcomes

Description

Utility wrapper function for computing out-of-sample metrics for a BART model when the test set outcomes are known.

Usage

```
bart_predict_for_test_data(bart_machine, Xtest, ytest)
```

Arguments

<code>bart_machine</code>	An object of class “bartMachine”.
<code>Xtest</code>	Data frame for test data containing rows at which predictions are to be made. Colnames should match that of the training data.
<code>ytest</code>	Actual outcomes for test data.

Value

For regression models, a list with the following components is returned:

y_hat	Predictions (as posterior means) for the test observations.
L1_err	L1 error for predictions.
L2_err	L2 error for predictions.
rmse	RMSE for predictions.

For classification models, a list with the following components is returned:

y_hat	Class predictions for the test observations.
confusion_matrix	A confusion matrix for the test observations.

Author(s)

Adam Kapelner and Justin Bleich

See Also

[predict](#)

Examples

```
#generate Friedman data
set.seed(11)
n = 250
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##split into train and test
train_X = X[1 : 200, ]
test_X = X[201 : 250, ]
train_y = y[1 : 200]
test_y = y[201 : 250]

##build BART regression model
bart_machine = build_bart_machine(train_X, train_y)

#explore performance on test data
oos_perf = bart_predict_for_test_data(bart_machine, test_X, test_y)
print(oos_perf$rmse)

#destroy BART model
destroy_bart_machine(bart_machine)
```

build_bart_machine	<i>Build a BART Model</i>
--------------------	---------------------------

Description

Builds a BART model for regression or classification.

Usage

```
build_bart_machine(X = NULL, y = NULL, Xy = NULL,
  num_trees = 50,
  num_burn_in = 250,
  num_iterations_after_burn_in = 1000,
  alpha = 0.95, beta = 2, k = 2, q = 0.9, nu = 3,
  prob_rule_class = 0.5,
  mh_prob_steps = c(2.5, 2.5, 4)/9,
  debug_log = FALSE,
  run_in_sample = TRUE,
  s_sq_y = "mse",
  cov_prior_vec = NULL,
  use_missing_data = FALSE,
  covariates_to_permute = NULL,
  num_rand_samps_in_library = 10000,
  use_missing_data_dummies_as_covars = FALSE,
  replace_missing_data_with_x_j_bar = FALSE,
  impute_missingness_with_rf_impute = FALSE,
  impute_missingness_with_x_j_bar_for_lm = TRUE,
  mem_cache_for_speed = TRUE,
  verbose = TRUE)
```

Arguments

X	Data frame of predictors. Factors are automatically converted to dummies internally.
y	Vector of response variable. If y is numeric or integer, a BART model for regression is built. If y is a factor with two levels, a BART model for classification is built.
Xy	A data frame of predictors and the response. The response column must be named "y".
num_trees	The number of trees to be grown in the sum-of-trees model.
num_burn_in	Number of MCMC samples to be discarded as "burn-in".
num_iterations_after_burn_in	Number of MCMC samples to draw from the posterior distribution of $\hat{f}(x)$.
alpha	Base hyperparameter in tree prior for whether a node is nonterminal or not.
beta	Power hyperparameter in tree prior for whether a node is nonterminal or not.
k	For regression, k determines the prior probability that $E(Y X)$ is contained in the interval (y_{min}, y_{max}) , based on a normal distribution. For example, when $k = 2$, the prior probability is 95%. For classification, k determines the prior probability that $E(Y X)$ is between $(-3, 3)$. Note that a larger value of k results in more shrinkage and a more conservative fit.

q	Quantile of the prior on the error variance at which the data-based estimate is placed. Note that the larger the value of q, the more aggressive the fit as you are placing more prior weight on values lower than the data-based estimate. Not used for classification.
nu	Degrees of freedom for the inverse χ^2 prior. Not used for classification.
prob_rule_class	Threshold for classification. Any observation with a conditional probability greater than prob_class_rule is assigned the “positive” outcome. Note that the first level of the response is treated as the “negative” outcome and the second is treated as the “positive” outcome.
mh_prob_steps	Vector of prior probabilities for proposing changes to the tree structures: (GROW, PRUNE, CHANGE)
debug_log	If TRUE, additional information about the model construction are printed to a file in the working directory.
run_in_sample	If TRUE, in-sample statistics such as $\hat{f}(x)$, Pseudo- R^2 , and RMSE are computed. Setting this to FALSE when not needed can decrease computation time.
s_sq_y	If “mse”, a data-based estimated of the error variance is computed as the MSE from ordinary least squares regression. If “var”, the data-based estimate is computed as the variance of the response. Not used in classification.
cov_prior_vec	Vector assigning relative weights to how often a particular variable should be proposed as a candidate for a split. The vector is internally normalized so that the weights sum to 1. Note that the length of this vector must equal the length of the design matrix after dummification and augmentation of indicators of missingness (if used). To see what the dummified matrix looks like, use dummify_data . See Bleich et al. (2013) for more details on when this feature is most appropriate.
use_missing_data	If TRUE, the missing data feature is used to automatically handle missing data without imputation. See Kapelner and Bleich (2013) for details.
covariates_to_permute	Private argument for cov_importance_test . Not needed by user.
num_rand_samps_in_library	Before building a BART model, samples from the Standard Normal and $\chi^2(\nu)$ are drawn to be used in the MCMC steps. This parameter determines the number of samples to be taken.
use_missing_data_dummies_as_covars	If TRUE, additional indicator variables for whether or not an observation in a particular column is missing are included. See Kapelner and Bleich (2013) for details.
replace_missing_data_with_x_j_bar	If TRUE, missing entries in X are imputed with average value or modal category.
impute_missingness_with_rf_impute	If TRUE, missing entries are filled in using the <code>rf.impute()</code> function from the <code>randomForest</code> library.
impute_missingness_with_x_j_bar_for_lm	If TRUE, when computing the data-based estimate of σ^2 , missing entries are imputed with average value or modal category.
mem_cache_for_speed	Speed enhancement that caches the predictors and the split values that are available at each node for selecting new rules. If the number of predictors is large, the

memory requirements become large. We recommend keeping this on (default) and turning it off if you experience out-of-memory errors.

`verbose` Prints information about progress of the algorithm to the screen.

Details

Returns an object of class “bartMachine”. Note that this object persists in the Java heap until `destroy_bart_machine` is called or the R session is terminated. The “bartMachine” object contains a list of the following components:

Value

<code>java_bart_machine</code>	A pointer to the BART Java object.
<code>train_data_features</code>	The names of the variables used in the training data.
<code>training_data_features_with_missing_features</code>	The names of the variables used in the training data. If <code>use_missing_data_dummies_as_covars = TRUE</code> this also includes dummies for any predictors that contain at least one missing entry (named “M_<feature>”).
<code>y</code>	The values of the response for the training data.
<code>y_levels</code>	The levels of the response (for classification only).
<code>pred_type</code>	Whether the model was build for regression or classification.
<code>model_matrix_training_data</code>	The training data with factors converted to dummies.
<code>num_cores</code>	The number of cores used to build the BART model.
<code>sig_sq_est</code>	The data-based estimate of σ^2 used to create the prior on the error variance for the BART model.
<code>time_to_build</code>	Total time to build the BART model.
<code>y_hat_train</code>	The posterior means of $\hat{f}(x)$ for each observation. Only returned if <code>run_in_sample = TRUE</code> .
<code>residuals</code>	The model residuals given by <code>y - y_hat_train</code> . Only returned if <code>run_in_sample = TRUE</code> .
<code>L1_err_train</code>	L1 error on the training set. Only returned if <code>run_in_sample = TRUE</code> .
<code>L2_err_train</code>	L2 error on the training set. Only returned if <code>run_in_sample = TRUE</code> .
<code>PseudoRsqr</code>	Calculated as $1 - \text{SSE} / \text{SST}$ where SSE is the sum of square errors in the training data and SST is the sample variance of the response times $n - 1$. Only returned if <code>run_in_sample = TRUE</code> .
<code>rmse_train</code>	Root mean square error on the training set. Only returned if <code>run_in_sample = TRUE</code> .

Additionally, the parameters passed to the function `build_bart_machine` are also components of the list.

Note

This function is parallelized by the number of cores set by `set_bart_machine_num_cores`. Each core will create an independent MCMC chain of size `num_burn_in + num_iterations_after_burn_in / bart_machine_num_cores`.

Author(s)

Adam Kapelner and Justin Bleich

References

HA Chipman, EI George, and RE McCulloch. BART: Bayesian Additive Regressive Trees. The Annals of Applied Statistics, 4(1): 266–298, 2010.

A Kapelner and J Bleich. Prediction with Missing Data via Bayesian Additive Regression Trees. ArXiv e-prints, 2013.

J Bleich, A Kapelner, ST Jensen, and EI George. Variable Selection Inference for Bayesian Additive Regression Trees. ArXiv e-prints, 2013.

See Also

[destroy_bart_machine](#), [build_bart_machine_cv](#)

Examples

```
##regression example

##generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)
summary(bart_machine)

#destroy BART model
destroy_bart_machine(bart_machine)

## Not run:
##Build another BART regression model
bart_machine = build_bart_machine(X,y, num_trees = 200, num_burn_in = 500,
num_iterations_after_burn_in = 1000)

#Destroy BART model
destroy_bart_machine(bart_machine)

##Classification example

#get data and only use 2 factors
data(iris)
iris2 = iris[51:150,]
iris2$Species = factor(iris2$Species)

#build BART classification model
bart_machine = build_bart_machine(iris2[,1:4], iris2$Species)

##get estimated probabilities
phat = bart_machine$p_hat_train
##look at in-sample confusion matrix
bart_machine$confusion_matrix

#destroy BART model
destroy_bart_machine(bart_machine)
```

```
## End(Not run)
```

```
build_bart_machine_cv Build BART-CV
```

Description

Builds a BART-CV model by cross-validating over a grid of hyperparameter choices.

Usage

```
build_bart_machine_cv(X = NULL, y = NULL, Xy = NULL,
  num_tree_cvs = c(50, 200), k_cvs = c(2, 3, 5),
  nu_q_cvs = list(c(3, 0.9), c(3, 0.99), c(10, 0.75)), k_folds = 5, ...)
```

Arguments

<code>X</code>	Data frame of predictors. Factors are automatically converted to dummies internally.
<code>y</code>	Vector of response variable. If <code>y</code> is numeric or integer, a BART model for regression is built. If <code>y</code> is a factor with two levels, a BART model for classification is built.
<code>Xy</code>	A data frame of predictors and the response. The response column must be named “y”.
<code>num_tree_cvs</code>	Vector of sizes for the sum-of-trees models to cross-validate over.
<code>k_cvs</code>	Vector of choices for the hyperparameter <code>k</code> to cross-validate over.
<code>nu_q_cvs</code>	List of vectors containing (nu, q) ordered pair choices to cross-validate over.
<code>k_folds</code>	Number of folds for cross-validation
<code>...</code>	Additional arguments to be passed to <code>build_bart_machine</code> .

Value

Returns an object of class “bartMachine” with the set of hyperparameters chosen via cross-validation.

Note

This function may require significant run-time. This function is parallelized by the number of cores set in `set_bart_machine_num_cores` via calling `build_bart_machine`.

Author(s)

Adam Kapelner and Justin Bleich

References

HA Chipman, EI George, and RE McCulloch. BART: Bayesian Additive Regressive Trees. The Annals of Applied Statistics, 4(1): 266–298, 2010.

See Also

[build_bart_machine](#)

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine_cv = build_bart_machine_cv(X, y)

#information about cross-validated model
summary(bart_machine_cv)

#destroy BART model
destroy_bart_machine(bart_machine_cv)

## End(Not run)
```

calc_credible_intervals

Calculate Credible Intervals

Description

Generates credible intervals for $\hat{f}(x)$ for a specified set of observations.

Usage

```
calc_credible_intervals(bart_machine, new_data,
ci_conf = 0.95)
```

Arguments

bart_machine	An object of class “bartMachine”.
new_data	A data frame containing observations at which credible intervals for $\hat{f}(x)$ are to be computed.
ci_conf	Confidence level for the credible intervals. The default is 95%.

Details

This interval is the appropriate quantiles based on the confidence level, `ci_conf`, of the predictions for each of the Gibbs samples post-burn in.

Value

Returns a matrix of the lower and upper bounds of the credible intervals for each observation in `new_data`.

Note

This function is parallelized by the number of cores set in `set_bart_machine_num_cores`.

Author(s)

Adam Kapelner and Justin Bleich

See Also

`calc_prediction_intervals`, `bart_machine_get_posterior`

Examples

```
#generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)

#get credible interval
cred_int = calc_credible_intervals(bart_machine, X)
print(head(cred_int))

#destroy BART model
destroy_bart_machine(bart_machine)
```

`calc_prediction_intervals`

Calculate Prediction Intervals

Description

Generates prediction intervals for \hat{y} for a specified set of observations.

Usage

```
calc_prediction_intervals(bart_machine, new_data,
  pi_conf = 0.95, normal_samples_per_gibbs_sample = 100)
```

Arguments

bart_machine	An object of class “bartMachine”.
new_data	A data frame containing observations at which prediction intervals for \hat{y} are to be computed.
pi_conf	Confidence level for the prediction intervals. The default is 95%.
normal_samples_per_gibbs_sample	The number of samples taken from the putative error distribution for each Gibbs sample after burn-in (see details). The default is 100.

Details

Credible intervals (see [calc_credible_intervals](#)) are the appropriate quantiles of the prediction for each of the Gibbs samples post-burn in. Prediction intervals also make use of the noise estimate at each Gibbs sample and hence are wider. For each Gibbs sample, we record the \hat{y} estimate of the response and the $\hat{\sigma}^2$ estimate of the noise variance. We then sample `normal_samples_per_gibbs_sample` times from a $N(\hat{y}, \hat{\sigma}^2)$ random variable to simulate many possible disturbances for that Gibbs sample. Then, all `normal_samples_per_gibbs_sample` times the number of Gibbs sample post burn-in are collected and the appropriate quantiles are taken based on the confidence level, `pi_conf`.

Value

Returns a matrix of the lower and upper bounds of the prediction intervals for each observation in `new_data`.

Note

This function is parallelized by the number of cores set in [set_bart_machine_num_cores](#).

Author(s)

Adam Kapelner and Justin Bleich

References

Kapelner, A and Bleich, J. bartMachine: A Powerful Tool for Machine Learning in R, arXiv preprints, 2013

See Also

[calc_credible_intervals](#), [bart_machine_get_posterior](#)

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)
```

```
#get prediction interval
pred_int = calc_prediction_intervals(bart_machine, X)
print(head(pred_int))

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

check_bart_error_assumptions

Check BART Error Assumptions

Description

Diagnostic tools to assess whether the errors of the BART model for regression are normally distributed and homoskedastic, as assumed by the model. This function generates a normal quantile plot of the residuals with a Shapiro-Wilks p-value as well as a residual plot.

Usage

```
check_bart_error_assumptions(bart_machine, hetero_plot = "yhats")
```

Arguments

bart_machine	An object of class “bartMachine”.
hetero_plot	If “yhats”, the residuals are plotted against the fitted values of the response. If “ys”, the residuals are plotted against the actual values of the response.

Value

None.

Author(s)

Adam Kapelner and Justin Bleich

See Also

[plot_convergence_diagnostics](#)

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 300
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)
```

```

##build BART regression model
bart_machine = build_bart_machine(X, y)

#check error diagnostics
check_bart_error_assumptions(bart_machine)

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)

```

cov_importance_test	<i>Importance Test for Covariate(s) of Interest</i>
---------------------	---

Description

This function tests the null hypothesis H_0 : These covariates of interest do not affect the response under the assumptions of the BART model.

Usage

```

cov_importance_test(bart_machine, covariates = NULL,
  num_permutation_samples = 100, plot = TRUE)

```

Arguments

bart_machine	An object of class “bart_machine”.
covariates	A vector of names of covariates of interest to be tested for having an effect on the response. A value of NULL indicates an omnibus test for all covariates having an effect on the response. If the name of a covariate is a factor, the entire factor will be permuted. We do not recommend entering the names of factor covariate dummies.
num_permutation_samples	The number of times to permute the covariates of interest and create a corresponding new BART model (see details).
plot	If TRUE, this produces a histogram of the Pseudo-Rsq's / total misclassification error rates from the num_permutations BART models created with the covariates permuted. The plot also illustrates the observed Pseudo-Rsq's / total misclassification error rate from the original training data and indicates the test's p-value.

Details

To test the importance of a covariate or a set of covariates of interest on the response, this function generates num_permutations BART models with the covariate(s) of interest permuted (differently each time). On each run, a measure of fit is recorded. For regression, the metric is Pseudo-Rsq; for classification, it is total misclassification error.

A p-value can then be generated as follows. For regression, the p-value is the number of permutation-sampled Pseudo-Rsq's greater than the observed Pseudo-Rsq divided by num_permutations + 1. For classification, the p-value is the number of permutation-sampled total misclassification errors less than the observed total misclassification error divided by num_permutations + 1.

Value

permutation_samples_of_error	A vector which records the error metric of the BART models with the covariates permuted (see details).
observed_error_estimate	For regression, this is the Pseudo-Rsq on the original training data set. For classification, this is the observed total misclassification error on the original training data set.
pval	The approximate p-value for this test (see details).

Note

This function is parallelized by the number of cores set in [set_bart_machine_num_cores](#).

Author(s)

Adam Kapelner and Justin Bleich

References

A Kapelner and J Bleich. `bartMachine`: A Powerful Tool for Machine Learning in R, arXiv preprints, 2013

Examples

```
## Not run:
##regression example

##generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)

##now test if X[, 1] affects Y nonparametrically under the BART model assumptions
cov_importance_test(bart_machine, covariates = c(1))
## note the plot and the printed p-value

##destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

destroy_bart_machine	<i>Destroy BART Model in Java</i>
----------------------	-----------------------------------

Description

Destroys a BART model by setting all Java pointers to null.

Usage

```
destroy_bart_machine(bart_machine)
```

Arguments

bart_machine

Details

Removing a “bart_machine” object from R does not free heap space from Java. Since BART objects can consume a large amount of RAM, it is important to remove these objects by calling this function if they are no longer needed or many BART objects are being created.

Value

None.

Author(s)

Adam Kapelner and Justin Bleich

Examples

```
##Generate Friedman Data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model and destroy it
bart_machine = build_bart_machine(X, y)

##should be called when object is no longer needed
##and before potentially removing the object from R
destroy_bart_machine(bart_machine)
```

dummify_data	<i>Dummify Design Matrix</i>
--------------	------------------------------

Description

Create a data frame with factors converted to dummies.

Usage

```
dummify_data(data)
```

Arguments

data	Data frame to be dummified.
------	-----------------------------

Details

The column names of the dummy variables are given by the “FactorName_LevelName” and are augmented to the end of the design matrix. See the example below.

Value

Returns a data frame with factors converted to dummy indicator variables.

Note

BART handles dummification internally. This function is provided as a utility function.

Author(s)

Adam Kapelner and Justin Bleich

Examples

```
#generate data
set.seed(11)
x1 = rnorm(20)
x2 = as.factor(ifelse(x1 > 0, "A", "B"))
x3 = runif(20)
X = data.frame(x1,x2,x3)
#dummify data
X_dummified = dummify_data(X)
print(X_dummified)
```

get_sigsqs

*Get Posterior Error Variance Estimates***Description**

Returns the posterior estimates of the error variance from the Gibbs samples with an option to create a histogram of the posterior estimates of the error variance with a credible interval overlaid.

Usage

```
get_sigsqs(bart_machine, after_burn_in = T,
plot_hist = F, plot_CI = .95, plot_sigma = F)
```

Arguments

bart_machine	An object of class “bartMachine”.
after_burn_in	If TRUE, only the σ^2 draws after the burn-in period are returned.
plot_hist	If TRUE, a histogram of the posterior σ^2 draws is generated.
plot_CI	Confidence level for credible interval on histogram.
plot_sigma	If TRUE, plots σ instead of σ^2 .

Value

Returns a vector of posterior σ^2 draws (with or without the burn-in samples).

Author(s)

Adam Kapelner and Justin Bleich

See Also

[get_sigsqs](#)

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 300
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)

#get posterior sigma^2's after burn-in and plot
sigsqs = get_sigsqs(bart_machine, plot_hist = TRUE)

#destroy BART model
destroy_bart_machine(bart_machine)
```

```
## End(Not run)
```

```
get_var_counts_over_chain
```

Get the Variable Inclusion Counts

Description

Computes the variable inclusion counts for a BART model.

Usage

```
get_var_counts_over_chain(bart_machine, type = "splits")
```

Arguments

bart_machine	An object of class “bartMachine”.
type	If “splits”, then the number of times each variable is chosen for a splitting rule is computed. If “trees”, then the number of times each variable appears in a tree is computed.

Value

Returns a matrix of counts of each predictor across all trees by Gibbs sample. Thus, the dimension is num_interactions_after_burn_in by p (where p is the number of predictors after dummifying factors and adding missingness dummies if specified by use_missing_data_dummies_as_covars).

Author(s)

Adam Kapelner and Justin Bleich

See Also

[get_var_props_over_chain](#)

Examples

```
## Not run:

#generate Friedman data
set.seed(11)
n = 200
p = 10
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y, num_trees = 20)

#get variable inclusion counts
var_counts = get_var_counts_over_chain(bart_machine)
print(var_counts)
```

```
#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

```
get_var_props_over_chain
```

Get the Variable Inclusion Proportions

Description

Computes the variable inclusion proportions for a BART model.

Usage

```
get_var_props_over_chain(bart_machine, type = "splits")
```

Arguments

bart_machine	An object of class “bartMachine”.
type	If “splits”, then the proportion of times each variable is chosen for a splitting rule versus all splitting rules is computed. If “trees”, then the proportion of times each variable appears in a tree versus all appearances of variables in trees is computed.

Value

Returns a vector of the variable inclusion proportions.

Author(s)

Adam Kapelner and Justin Bleich

See Also

[get_var_counts_over_chain](#)

Examples

```
#generate Friedman data
set.seed(11)
n = 200
p = 10
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y, num_trees = 20)

#Get variable inclusion proportions
var_props = get_var_props_over_chain(bart_machine)
print(var_props)
```

```
#destroy BART model  
destroy_bart_machine(bart_machine)
```

```
init_java_for_bart_machine_with_mem_in_mb
```

Initialize a JVM with a pre-specified heap size

Description

Initialize a Java virtual Machine (JVM) with a pre-specified heap size in MB.

Usage

```
init_java_for_bart_machine_with_mem_in_mb(bart_max_mem)
```

Arguments

bart_max_mem Java Virtual Machine heap size in MB

Details

This function can only be called once per R session and must be called before [build_bart_machine](#), else the default heap size of 1.5GB will be used.

Value

None.

Author(s)

Adam Kapelner and Justin Bleich

Examples

```
## Not run:  
##initialize a Java Virtual Machine with heap size of 3000MB  
##this should be run before any BART models are built  
init_java_for_bart_machine_with_mem_in_mb(3000)  
  
## End(Not run)
```

interaction_investigator

Explore Pairwise Interactions in BART Model

Description

Explore the pairwise interaction counts for a BART model to learn about interactions fit by the model. This function includes an option to generate a plot of the pairwise interaction counts.

Usage

```
interaction_investigator(bart_machine, plot = TRUE,
  num_replicates_for_avg = 5, num_trees_bottleneck = 20,
  num_var_plot = 50, cut_bottom = NULL, bottom_margin = 10)
```

Arguments

- | | |
|------------------------|---|
| bart_machine | An object of class “bartMachine”. |
| plot | If TRUE, a plot of the pairwise interaction counts is generated. |
| num_replicates_for_avg | The number of replicates of BART to be used to generate pairwise interaction inclusion counts. Averaging across multiple BART models improves stability of the estimates. |
| num_trees_bottleneck | Number of trees to be used in the sum-of-trees model for computing pairwise interactions counts. A small number of trees should be used to force the variables to compete for entry into the model. |
| num_var_plot | Number of variables to be shown on the plot. If “Inf,” all variables are plotted (not recommended if the number of predictors is large). Default is 50. |
| cut_bottom | A display parameter between 0 and 1 that controls where the y-axis is plotted. A value of 0 would begin the y-axis at 0; a value of 1 begins the y-axis at the minimum of the average pairwise interaction inclusion count (the smallest bar in the bar plot). Values between 0 and 1 begin the y-axis as a percentage of that minimum. |
| bottom_margin | A display parameter that adjusts the bottom margin of the graph if labels are clipped. The scale of this parameter is the same as set with <code>par(mar = c(...))</code> in R. Higher values allow for more space if the crossed covariate names are long. Note that making this parameter too large will prevent plotting and the plot function in R will throw an error. |

Details

An interaction between two variables is considered to occur whenever a path from any node of a tree to any of its terminal node contains splits using those two variables. See Kapelner and Bleich, 2013, Section 4.11.

Value

interaction_counts_avg

For each of the $p \times p$ interactions, what is the average count across all num_replicates_for_avg BART model replicates' post burn-in Gibbs samples in all trees.

interaction_counts_sd

For each of the $p \times p$ interactions, what is the average sd of the interaction counts across the num_replicates_for_avg BART models replicates.

Note

In the plot, the red bars correspond to the standard error of the variable inclusion proportion estimates (since multiple replicates were used).

Author(s)

Adam Kapelner and Justin Bleich

References

A Kapelner and J Bleich. bartMachine: A Powerful Tool for Machine Learning in R, arXiv preprints, 2013

See Also

[investigate_var_importance](#)

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 200
p = 10
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y, num_trees = 20)

#investigate interactions
interaction_investigator(bart_machine)

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

investigate_var_importance

Explore Variable Inclusion Proportions in BART Model

Description

Explore the variable inclusion proportions for a BART model to learn about the relative influence of the different covariates. This function includes an option to generate a plot of the variable inclusion proportions.

Usage

```
investigate_var_importance(bart_machine, type = "splits",
  plot = TRUE, num_replicates_for_avg = 5, num_trees_bottleneck = 20,
  num_var_plot = Inf, bottom_margin = 10)
```

Arguments

<code>bart_machine</code>	An object of class “bartMachine”.
<code>type</code>	If “splits”, then the proportion of times each variable is chosen for a splitting rule is computed. If “trees”, then the proportion of times each variable appears in a tree is computed.
<code>plot</code>	If TRUE, a plot of the variable inclusion proportions is generated.
<code>num_replicates_for_avg</code>	The number of replicates of BART to be used to generate variable inclusion proportions. Averaging across multiple BART models improves stability of the estimates. See Bleich et al. (2013) for more details.
<code>num_trees_bottleneck</code>	Number of trees to be used in the sum-of-trees for computing the variable inclusion proportions. A small number of trees should be used to force the variables to compete for entry into the model. Chipman et al. (2010) recommend 20. See this reference for more details.
<code>num_var_plot</code>	Number of variables to be shown on the plot. If “Inf”, all variables are plotted.
<code>bottom_margin</code>	A display parameter that adjusts the bottom margin of the graph if labels are clipped. The scale of this parameter is the same as set with <code>par(mar = c(...))</code> in R. Higher values allow for more space if the covariate names are long. Note that making this parameter too large will prevent plotting and the plot function in R will throw an error.

Details

In the plot, the red bars correspond to the standard error of the variable inclusion proportion estimates.

Value

Invisibly, returns a list with the following components:

<code>avg_var_props</code>	The average variable inclusion proportions for each variable (across <code>num_replicates_for_avg</code>)
----------------------------	--

sd_var_props The standard deviation of the variable inclusion proportions for each variable (across num_replicates_for_avg)

Note

This function is parallelized by the number of cores set in [set_bart_machine_num_cores](#).

Author(s)

Adam Kapelner and Justin Bleich

References

J Bleich, A Kapelner, ST Jensen, and EI George. Variable Selection Inference for Bayesian Additive Regression Trees. ArXiv e-prints, 2013.

HA Chipman, EI George, and RE McCulloch. BART: Bayesian Additive Regressive Trees. The Annals of Applied Statistics, 4(1): 266–298, 2010.

See Also

[interaction_investigator](#)

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 200
p = 10
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y, num_trees = 20)

#investigate variable inclusion proportions
investigate_var_importance(bart_machine)

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

k_fold_cv

Estimate Out-of-sample Error with K-fold Cross validation

Description

Builds a BART model using a specified set of arguments to `build_bart_machine` and estimates the out-of-sample performance by using k-fold cross validation.

Usage

```
k_fold_cv(X, y, k_folds = 5, ...)
```

Arguments

X	Data frame of predictors. Factors are automatically converted to dummies internally.
y	Vector of response variable. If y is numeric or integer, a BART model for regression is built. If y is a factor with two levels, a BART model for classification is built.
k_folds	Number of folds to cross-validate over.
...	Additional arguments to be passed to <code>build_bart_machine</code> .

Details

For each fold, a new BART model is trained (using the same set of arguments) and its performance is evaluated on the holdout piece of that fold.

Value

For regression models, a list with the following components is returned:

L1_err	Aggregate L1 error across the folds.
L2_err	Aggregate L1 error across the folds.
rmse	Aggregate RMSE across the folds.

For classification models, a list with the following components is returned:

y_hat	Class predictions for the test observations.
confusion_matrix	Aggregate confusion matrix across the folds.
misclassification_error	Total misclassification error across the folds.

Note

This function is parallelized by the number of cores set in [set_bart_machine_num_cores](#).

Author(s)

Adam Kapelner and Justin Bleich

See Also

[build_bart_machine](#)

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

#evaluate default BART on 5 folds
k_fold_val = k_fold_cv(X, y)
print(k_fold_val$rmse)

## End(Not run)
```

pd_plot

Partial Dependence Plot

Description

Creates a partial dependence plot for a BART model for regression or classification.

Usage

```
pd_plot(bart_machine, j,
  levs = c(0.05, seq(from = 0.1, to = 0.9, by = 0.1), 0.95),
  lower_ci = 0.025, upper_ci = 0.975)
```

Arguments

bart_machine	An object of class “bartMachine”.
j	The number or name of the column in the design matrix for which the partial dependence plot is to be created.
levs	Quantiles at which the partial dependence function should be evaluated. Linear extrapolation is performed between these points.
lower_ci	Lower limit for credible interval
upper_ci	Upper limit for credible interval

Details

For regression models, the units on the y-axis are the same as the units of the response. For classification models, the units on the y-axis are probits.

Value

Invisibly, returns a list with the following components:

x_j_quants	Quantiles at which the partial dependence function is evaluated.
bart_avg_predictions_by_quantile	Posterior means for $\hat{f}(x)$ at x_j_quants.

Note

This function is parallelized by the number of cores set in `set_bart_machine_num_cores`.

Author(s)

Adam Kapelner and Justin Bleich

References

HA Chipman, EI George, and RE McCulloch. BART: Bayesian Additive Regressive Trees. The Annals of Applied Statistics, 4(1): 266–298, 2010.

Examples

```
## Not run:
#Regression example

#generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)

#partial dependence plot for quadratic term
pd_plot(bart_machine, "X3")

#destroy BART model
destroy_bart_machine(bart_machine)

#Classification example

#get data and only use 2 factors
data(iris)
iris2 = iris[51:150,]
iris2$Species = factor(iris2$Species)

#build BART classification model
bart_machine = build_bart_machine(iris2[,1:4], iris2$Species)

#partial dependence plot
pd_plot(bart_machine, "Petal.Width")

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

plot_convergence_diagnostics

Plot Convergence Diagnostics

Description

A suite of plots to assess convergence diagnostics and features of the BART model.

Usage

```
plot_convergence_diagnostics(bart_machine,
plots = c("sigsqs", "mh_acceptance", "num_nodes", "tree_depths"))
```

Arguments

bart_machine	An object of class “bartMachine”.
plots	The list of plots to be displayed. The four options are: "sigsqs", "mh_acceptance", "num_nodes", "tree_depths".

Details

The “sigsqs” option plots the posterior error variance estimates by the Gibbs sample number. This is a standard tool to assess convergence of MCMC algorithms. This option is not applicable to classification BART models.

The “mh_acceptance” option plots the proportion of Metropolis-Hastings steps accepted for each Gibbs sample (number accepted divided by number of trees).

The “num_nodes” option plots the average number of nodes across each tree in the sum-of-trees model by the Gibbs sample number (for post burn-in only). The blue line is the average number of nodes over all trees.

The “tree_depths” option plots the average tree depth across each tree in the sum-of-trees model by the Gibbs sample number (for post burn-in only). The blue line is the average number of nodes over all trees.

Value

None.

Note

The “sigsqs” plot separates the burn-in σ^2 's for the first core by post burn-in σ^2 's estimates for all cores by grey vertical lines. The “mh_acceptance” plot separates burn-in from post-burn in by a grey vertical line. Post burn-in, the different core proportions plot in different colors. The “num_nodes” plot separates different core estimates by vertical lines (post burn-in only). The “tree_depths” plot separates different core estimates by vertical lines (post burn-in only).

Author(s)

Adam Kapelner and Justin Bleich

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)

#plot convergence diagnostics
plot_convergence_diagnostics(bart_machine)

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

plot_y_vs_yhat

Plot the fitted Versus Actual Response

Description

Generates a plot actual versus fitted values and corresponding credible intervals or prediction intervals for the fitted values.

Usage

```
plot_y_vs_yhat(bart_machine, Xtest = NULL, ytest = NULL,
  credible_intervals = FALSE, prediction_intervals = FALSE,
  interval_confidence_level = 0.95)
```

Arguments

bart_machine	An object of class “bartMachine”.
Xtest	Optional argument for test data. If included, BART computes fitted values at the rows of Xtest. Else, the fitted values from the training data are used.
ytest	Optional argument for test data. Vector of observed values corresponding to the rows of Xtest to be plotted against the predictions for the rows of Xtest.
credible_intervals	If TRUE, Bayesian credible intervals are computed using the quantiles of the posterior distribution of $\hat{f}(x)$. See calc_credible_intervals for details.
prediction_intervals	If TRUE, Bayesian predictive intervals are computed using the a draw of from $\hat{f}(x)$. See calc_prediction_intervals for details.
interval_confidence_level	Desired level of confidence for credible or prediction intervals.

Value

None.

Note

This function is parallelized by the number of cores set in [set_bart_machine_num_cores](#).

Author(s)

Adam Kapelner and Justin Bleich

See Also

[bart_machine_get_posterior](#), [calc_credible_intervals](#), [calc_prediction_intervals](#)

Examples

```
## Not run:
#generate linear data
set.seed(11)
n = 500
p = 3
X = data.frame(matrix(runif(n * p), ncol = p))
y = 3*X[,1] + 2*X[,2] + X[,3] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)

##generate plot
plot_y_vs_yhat(bart_machine)

#generate plot with prediction bands
plot_y_vs_yhat(bart_machine, prediction_intervals = TRUE)

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

predict.bartMachine *Make a prediction on data using a BART object*

Description

Makes a prediction on new data given a fitted BART model for regression or classification.

Usage

```
## S3 method for class 'bartMachine'
predict(object, new_data, type, prob_rule_class, ...)
```


Arguments

object	An object of class “bartMachine”.
new_data	A data frame where each row is an observation to predict. The column names should be the same as the column names of the training data.
type	Only relevant if the bartMachine model is classification. The type can be “prob” which will return the estimate of $P(Y = 1)$ (the “positive” class) or “class” which will return the best guess as to the class of the object, in the original label, based on if the probability estimate is greater than prob_rule_class. Default is “prob.”
prob_rule_class	The rule to determine when the class estimate is $Y = 1$ (the “positive” class) based on the probability estimate. This defaults to what was originally specified in the bart_machine object.
...	Parameters that are ignored.

Value

If regression, a numeric vector of \hat{y} , the best guess as to the response. If classification and `type = “prob”`, a numeric vector of \hat{p} , the best guess as to the probability of the response class being the “positive” class. If classification and `type = “class”`, a character vector of the best guess of the response’s class labels.

Author(s)

Adam Kapelner and Justin Bleich

Examples

```
#Regression example
## Not run:
#generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)

##make predictions on the training data
y_hat = predict(bart_machine, X)

##destroy BART model
destroy_bart_machine(bart_machine)

#Classification example
data(iris)
iris2 = iris[51 : 150, ] #do not include the third type of flower for this example
iris2$Species = factor(iris2$Species)
bart_machine = build_bart_machine(iris2[,1:4], iris2$Species)

##make probability predictions on the training data
p_hat = predict(bart_machine, X)
```

```
##make class predictions on test data
y_hat_class = predict(bart_machine, X, type = "class")

##make class predictions on test data conservatively for 'versicolor'
y_hat_class_conservative = predict(bart_machine, X, type = "class", prob_rule_class = 0.9)

##destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

print.bartMachine	<i>Summarizes information about a bartMachine object.</i>
-------------------	---

Description

This is an alias for the [summary](#) function. See description in that section.

Usage

```
## S3 method for class 'bartMachine'
print(x, ...)
```

Arguments

x	An object of class “bartMachine”.
...	Parameters that are ignored.

Value

None.

Author(s)

Adam Kapelner and Justin Bleich

Examples

```
#Regression example

#generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)

##print out details
```

```
print(bart_machine)

##Also, the default print works too
bart_machine

##destroy BART model
destroy_bart_machine(bart_machine)
```

rmse_by_num_trees	<i>Assess the Out-of-sample RMSE by Number of Trees</i>
-------------------	---

Description

Assess out-of-sample RMSE of a BART model for varying numbers of trees in the sum-of-trees model.

Usage

```
rmse_by_num_trees(bart_machine, tree_list = c(5, seq(10, 50, 10), 100, 150, 200),
  in_sample = FALSE, plot = TRUE, holdout_pctg = 0.3, num_replicates = 4)
```

Arguments

<code>bart_machine</code>	An object of class “bartMachine”.
<code>tree_list</code>	List of sizes for the sum-of-trees models.
<code>in_sample</code>	If TRUE, the RMSE is computed on in-sample data rather than an out-of-sample holdout.
<code>plot</code>	If TRUE, a plot of the RMSE by the number of trees in the ensemble is created.
<code>holdout_pctg</code>	Percentage of the data to be treated as an out-of-sample holdout.
<code>num_replicates</code>	Number of replicates to average the results over. Each replicate uses a randomly sampled holdout of the data, (which could have overlap).

Value

Invisibly, returns the out-of-sample average RMSEs for each tree size.

Note

Since using a large number of trees can substantially increase computation time, this plot can help assess whether a smaller ensemble size is sufficient to obtain desirable predictive performance. This function is parallelized by the number of cores set in [set_bart_machine_num_cores](#).

Author(s)

Adam Kapelner and Justin Bleich

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 200
p = 10
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y, num_trees = 20)

#explore RMSE by number of trees
rmse_by_num_trees(bart_machine)

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

```
set_bart_machine_num_cores
```

Set the Number of Cores for BART

Description

Sets the number of cores to be used for all parallelized BART functions.

Usage

```
set_bart_machine_num_cores(num_cores)
```

Arguments

num_cores Number of cores to use

Value

None.

Author(s)

Adam Kapelner and Justin Bleich

See Also

[bart_machine_num_cores](#)

Examples

```
## Not run:
## set all parallelized functions to use 4 cores
## set_bart_machine_num_cores(4)

## End(Not run)
```

summary.bartMachine	<i>Summarizes information about a bartMachine object.</i>
---------------------	---

Description

Provides a quick summary of the BART model.

Usage

```
## S3 method for class 'bartMachine'
summary(object, ...)
```

Arguments

object	An object of class “bartMachine”.
...	Parameters that are ignored.

Details

Gives the version number of the bartMachine package used to build this bartMachine object and if the object models either “regression” or “classification.” Gives the amount of training data and the dimension of feature space. Prints the amount of time it took to build the model, how many processor cores were used to during its construction, as well as the number of burn-in and posterior Gibbs samples were used.

If the model is for regression, it prints the estimate of σ^2 before the model was constructed as well as after so the user can inspect how much variance was explained.

If the model was built using the run_in_sample = TRUE parameter in [build_bart_machine](#) and is for regression, the summary L1, L2, rmse, Pseudo- R^2 are printed as well as the p-value for the tests of normality and zero-mean noise. If the model is for classification, a confusion matrix is printed.

Value

None.

Author(s)

Adam Kapelner and Justin Bleich

Examples

```
#Regression example

#generate Friedman data
set.seed(11)
n = 200
p = 5
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model
bart_machine = build_bart_machine(X, y)

##print out details
summary(bart_machine)

##Also, the default print works too
bart_machine

##destroy BART model
destroy_bart_machine(bart_machine)
```

var_selection_by_permute_response_cv

Perform Variable Selection Using Cross-validation Procedure

Description

Performs variable selection by cross-validating over the three threshold-based procedures outlined in Bleich et al. (2013) and selecting the single procedure that returns the lowest cross-validation RMSE.

Usage

```
var_selection_by_permute_response_cv(bart_machine, k_folds = 5,
num_reps_for_avg = 5, num_permute_samples = 100,
num_trees_for_permute = 20, alpha = 0.05, num_trees_pred_cv = 50)
```

Arguments

bart_machine	An object of class “bartMachine”.
k_folds	Number of folds to be used in cross-validation.
num_reps_for_avg	Number of replicates to over over to for the BART model’s variable inclusion proportions.
num_permute_samples	Number of permutations of the response to be made to generate the “null” permutation distribution.
num_trees_for_permute	Number of trees to use in the variable selection procedure. As with investigate_var_importance , a small number of trees should be used to force variables to compete for entry into the model. Note that this number is used to estimate both the “true” and “null” variable inclusion proportions.

alpha Cut-off level for the thresholds.

num_trees_pred_cv Number of trees to use for prediction on the hold-out portion of each fold. Once variables have been selected using the training portion of each fold, a new model is built using only those variables with num_trees_pred_cv trees in the sum-of-trees model. Forecasts for the holdout sample are made using this model. A larger number of trees is recommended to exploit the full forecasting power of BART.

Details

See Bleich et al. (2013) for a complete description of the procedures outlined above as well as the corresponding vignette for a brief summary with examples.

Value

Returns a list with the following components:

best_method The name of the best variable selection procedure, as chosen via cross-validation.

important_vars_cv The variables chosen by the best_method above.

Note

This function can have substantial run-time. This function is parallelized by the number of cores set in [set_bart_machine_num_cores](#).

Author(s)

Adam Kapelner and Justin Bleich

References

J Bleich, A Kapelner, ST Jensen, and EI George. Variable Selection Inference for Bayesian Additive Regression Trees. ArXiv e-prints, 2013.

A Kapelner and J Bleich. bartMachine: A Powerful Tool for Machine Learning in R, arXiv preprints, 2013

See Also

[var_selection_by_permute_response_three_methods](#), [investigate_var_importance](#)

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 150
p = 100 ##95 useless predictors
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model (not actually used in variable selection)
bart_machine = build_bart_machine(X, y)
```

```
#variable selection via cross-validation
var_sel_cv = var_selection_by_permute_response_cv(bart_machine, k_folds = 3)
print(var_sel_cv$best_method)
print(var_sel_cv$important_vars_cv)

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

var_selection_by_permute_response_three_methods

Perform Variable Selection using Three Threshold-based Procedures

Description

Performs variable selection using the three thresholding methods introduced in Bleich et al. (2013).

Usage

```
var_selection_by_permute_response_three_methods(bart_machine,
num_reps_for_avg = 10, num_permute_samples = 100,
num_trees_for_permute = 20, alpha = 0.05,
plot = TRUE, num_var_plot = Inf, bottom_margin = 10)
```

Arguments

bart_machine	An object of class “bartMachine”.
num_reps_for_avg	Number of replicates to over over to for the BART model’s variable inclusion proportions.
num_permute_samples	Number of permutations of the response to be made to generate the “null” permutation distribution.
num_trees_for_permute	Number of trees to use in the variable selection procedure. As with investigate_var_importance , a small number of trees should be used to force variables to compete for entry into the model. Note that this number is used to estimate both the “true” and “null” variable inclusion proportions.
alpha	Cut-off level for the thresholds.
plot	If TRUE, a plot showing which variables are selected by each of the procedures is generated.
num_var_plot	Number of variables (in order of decreasing variable inclusion proportion) to be plotted.
bottom_margin	A display parameter that adjusts the bottom margin of the graph if labels are clipped. The scale of this parameter is the same as set with <code>par(mar = c(...))</code> in R. Higher values allow for more space if the crossed covariate names are long. Note that making this parameter too large will prevent plotting and the plot function in R will throw an error.

Details

See Bleich et al. (2013) for a complete description of the procedures outlined above as well as the corresponding vignette for a brief summary with examples.

Value

Invisibly, returns a list with the following components:

important_vars_local_names	Names of the variables chosen by the Local procedure.
important_vars_global_max_names	Names of the variables chosen by the Global Max procedure.
important_vars_global_se_names	Names of the variables chosen by the Global SE procedure.
important_vars_local_col_nums	Column numbers of the variables chosen by the Local procedure.
important_vars_global_max_col_nums	Column numbers of the variables chosen by the Global Max procedure.
important_vars_global_se_col_nums	Column numbers of the variables chosen by the Global SE procedure.
var_true_props_avg	The variable inclusion proportions for the actual data.
permute_mat	The permutation distribution generated by permuting the response vector.

Note

Although the reference only explores regression settings, this procedure is applicable to both regression and classification problems. This function is parallelized by the number of cores set in [set_bart_machine_num_cores](#).

Author(s)

Adam Kapelner and Justin Bleich

References

J Bleich, A Kapelner, ST Jensen, and EI George. Variable Selection Inference for Bayesian Additive Regression Trees. ArXiv e-prints, 2013.

A Kapelner and J Bleich. bartMachine: A Powerful Tool for Machine Learning in R, arXiv preprints, 2013

See Also

[var_selection_by_permute_response_cv](#), [investigate_var_importance](#)

Examples

```
## Not run:
#generate Friedman data
set.seed(11)
n = 300
p = 20 ##15 useless predictors
```

```
X = data.frame(matrix(runif(n * p), ncol = p))
y = 10 * sin(pi * X[,1] * X[,2]) + 20 * (X[,3] - .5)^2 + 10 * X[,4] + 5 * X[,5] + rnorm(n)

##build BART regression model (not actually used in variable selection)
bart_machine = build_bart_machine(X, y)

#variable selection
var_sel = var_selection_by_permute_response_three_methods(bart_machine)
print(var_sel$important_vars_local_names)
print(var_sel$important_vars_global_max_names)

#destroy BART model
destroy_bart_machine(bart_machine)

## End(Not run)
```

Index

*Topic \textasciitildekw1

bart_machine_get_posterior, 2
bart_machine_num_cores, 3
bart_predict_for_test_data, 4
build_bart_machine, 6
build_bart_machine_cv, 10
calc_credible_intervals, 11
calc_prediction_intervals, 12
check_bart_error_assumptions, 14
cov_importance_test, 15
destroy_bart_machine, 17
dummify_data, 18
get_sigsqs, 19
get_var_counts_over_chain, 20
get_var_props_over_chain, 21
init_java_for_bart_machine_with_mem_in_mb, 22
interaction_investigator, 23
investigate_var_importance, 25
k_fold_cv, 26
pd_plot, 28
plot_convergence_diagnostics, 30
plot_y_vs_yhat, 31
predict.bartMachine, 32
print.bartMachine, 34
rmse_by_num_trees, 35
set_bart_machine_num_cores, 36
summary.bartMachine, 37
var_selection_by_permute_response_cv, 38
var_selection_by_permute_response_three_methods, 40
dummify_data, 18
get_sigsqs, 19
get_var_counts_over_chain, 20
get_var_props_over_chain, 21
init_java_for_bart_machine_with_mem_in_mb, 22
interaction_investigator, 23
investigate_var_importance, 25
k_fold_cv, 26
pd_plot, 28
plot_convergence_diagnostics, 30
plot_y_vs_yhat, 31
predict.bartMachine, 32
print.bartMachine, 34
rmse_by_num_trees, 35
set_bart_machine_num_cores, 36
summary.bartMachine, 37
var_selection_by_permute_response_cv, 38
var_selection_by_permute_response_three_methods, 40
bart_machine_get_posterior, 2, 12, 13, 32
bart_machine_num_cores, 3, 36
bart_predict_for_test_data, 4
build_bart_machine, 6, 10, 11, 22, 27, 37
build_bart_machine_cv, 9, 10
calc_credible_intervals, 3, 11, 13, 31, 32
calc_prediction_intervals, 3, 12, 12, 31, 32
check_bart_error_assumptions, 14
cov_importance_test, 7, 15
destroy_bart_machine, 8, 9, 17
dummify_data, 7, 18
get_sigsqs, 19, 19
get_var_counts_over_chain, 20, 21
get_var_props_over_chain, 20, 21
init_java_for_bart_machine_with_mem_in_mb, 22
interaction_investigator, 23, 26

*Topic \textasciitildekw2

bart_machine_get_posterior, 2
bart_machine_num_cores, 3
bart_predict_for_test_data, 4
build_bart_machine, 6
build_bart_machine_cv, 10
calc_credible_intervals, 11
calc_prediction_intervals, 12
check_bart_error_assumptions, 14
cov_importance_test, 15
destroy_bart_machine, 17

investigate_var_importance, [24](#), [25](#),
[38–41](#)

k_fold_cv, [26](#)

pd_plot, [28](#)

plot_convergence_diagnostics, [14](#), [30](#)

plot_y_vs_yhat, [31](#)

predict, [5](#)

predict.bartMachine, [32](#)

print.bartMachine, [34](#)

rmse_by_num_trees, [35](#)

set_bart_machine_num_cores, [2](#), [4](#), [8](#), [10](#),
[12](#), [13](#), [16](#), [26](#), [27](#), [29](#), [32](#), [35](#), [36](#), [39](#),
[41](#)

summary, [34](#)

summary.bartMachine, [37](#)

var_selection_by_permute_response_cv,
[38](#), [41](#)

var_selection_by_permute_response_three_methods,
[39](#), [40](#)